

Bachelorarbeit

Malleable TOHTN Planning using CrowdHTN and Mallob

Niko Wilhelm

Abgabedatum: 19.10.2012

Betreuer: Prof. Dr. Peter Sanders
M.Sc. Dominik Schreiber

Institut für Theoretische Informatik, Algorithmik
Fakultät für Informatik
Karlsruher Institut für Technologie

Hiermit versichere ich, dass ich diese Arbeit selbständig verfasst und keine anderen, als die angegebenen Quellen und Hilfsmittel benutzt, die wörtlich oder inhaltlich übernommenen Stellen als solche kenntlich gemacht und die Satzung des Karlsruher Instituts für Technologie zur Sicherung guter wissenschaftlicher Praxis in der jeweils gültigen Fassung beachtet habe.

Karlsruhe, den September 1, 2022

Zusammenfassung

Hier die deutsche Zusammenfassung.

Ich bin Blindtext. Von Geburt an. Es hat lange gedauert, bis ich begriffen habe, was es bedeutet, ein blinder Text zu sein: Man macht keinen Sinn. Man wirkt hier und da aus dem Zusammenhang gerissen. Oft wird man gar nicht erst gelesen. Aber bin ich deshalb ein schlechter Text? Ich weiß, dass ich nie die Chance haben werde im Stern zu erscheinen. Aber bin ich darum weniger wichtig? Ich bin blind! Aber ich bin gerne Text. Und sollten Sie mich jetzt tatsächlich zu Ende lesen, dann habe ich etwas geschafft, was den meisten „normalen“ Texten nicht gelingt.

Ich bin Blindtext. Von Geburt an. Es hat lange gedauert, bis ich begriffen habe, was es bedeutet, ein blinder Text zu sein: Man macht keinen Sinn. Man wirkt hier und da aus dem Zusammenhang gerissen. Oft wird man gar nicht erst gelesen. Aber bin ich deshalb ein schlechter Text? Ich weiß, dass ich nie die Chance haben werde im Stern zu erscheinen. Aber bin ich darum weniger wichtig? Ich bin blind! Aber ich bin gerne Text.

Abstract

And here an English translation of the German abstract.

I'm blind text. From birth. It took a long time until I realized what it means to be random text: You make no sense. You stand here and there out of context. Frequently, they do not even read. But I have a bad copy? I know that I will never have the chance of appearing in the. But I'm any less important? I'm blind! But I like to text. And you should see me now actually over, then I have accomplished something that is not possible in most "normal" copies.

I'm blind text. From birth. It took a long time until I realized what it means to be random text: You make no sense. You stand here and there out of context. Frequently, they do not even read. But I have a bad copy? I know that I will never have the chance of appearing in the. But I'm any less important? I'm blind! But I like to text.

Danksagungen

Thanks to i10pc135 which suffered much to make the experimental evaluation possible.

Inhaltsverzeichnis

1	Introduction	7
1.1	Motivation	7
1.2	Research Goal	7
1.3	Thesis Overview	7
2	Theoretical Foundations	7
2.1	TOHTN Formalism	8
2.1.1	Defining TOHTN Planning Problems	8
2.2	Malleable Algorithms	10
3	TOHTN Metadata	10
4	Implementation	10
5	Improvements to CrowdHTN	11
5.1	Reducing Memory Consumption	11
5.1.1	Efficiently Storing the Preceding Plan	11
5.1.2	Reducing Copies of the World State	11
5.1.3	Only Saving 'Potentially Interesting' Nodes	11
5.2	Making Crowd More CPU Efficient	11
5.3	Efficiently Hashing Nodes of the Search Graph	11
5.4	Preceding Plan	12
5.5	Lazy Instantiation of Child Nodes	12
5.6	Using Domain Meta-Information	13
5.7	Global and Memory Efficient Loop Detection	13
5.7.1	Restarts Under Loop Detection	14
5.7.2	Reaching Consensus on Search Version	14
5.7.3	Completeness Under Loop Detection	14
6	Theoretical Improvements to the Crowd Planner	15
6.1	Search Algorithms in Crowd	15
6.2	Loop Detection	15
6.2.1	Perfect Loop Detection	15
6.2.2	Bloom Filters	15
6.2.3	Distributed Loop Detection	15
7	Malleable TOHTN Planning	15
8	Experimental Evaluation	15
9	Future Work	15

Abbildungsverzeichnis

Tabellenverzeichnis

Algorithmenverzeichnis

1 Introduction

1.1 Motivation

1.2 Research Goal

- Provide a performant parallel TOHTN planner by improving upon the Crowd planner -
Provide integration of TOHTN into the Mallob malleable load balancer - Compare performance
of parallel to malleable TOHTN planning -

1.3 Thesis Overview

2 Theoretical Foundations

2.1 TOHTN Formalism

2.1.1 Defining TOHTN Planning Problems

Both HTN and TOHTN planning are based on the idea of decomposing a list of initial tasks down into smaller subtasks until those subtasks can be achieved by simple actions.

Multiple definitions for HTN planning exist. In this work we build on the definition introduced in [1].

Definition 1. A **predicate** consists of two parts. Firstly a predicate symbol $p \in \mathcal{P}$ where \mathcal{P} is the finite set of predicate symbols. Secondly of a list of terms τ_1, \dots, τ_k where each term τ_i is either a constant symbol $c \in \mathcal{C}$, with \mathcal{C} being the finite set of constant symbols, or a variable symbol $v \in \mathcal{V}$, where \mathcal{V} is the infinite set of variable symbols.

The set of all predicates is called \mathcal{Q} .

With the definition of a predicate in place, we can then define a grounding as well as our world state.

Definition 2. A **ground predicate** is a predicate where the terms contain no variable symbols or, in other words, a predicate that contains only constant symbols.

Definition 3. A **state** $s \in 2^{\mathcal{Q}}$ is a set of ground predicates for which we make the closed-world-assumption. Under the closed-world-assumption, only positive predicates are explicitly represented in s . All predicates not in s are implicitly negative.

Definition 4. With T_p the set of primitive task symbols, a **primitive task** t_p is defined as a triple $t_p(\tilde{t}_p(a_1, \dots, a_k), \text{pre}(t_p), \text{eff}(t_p))$. $\tilde{t}_p \in T_p$ is the task symbol, $a_1, \dots, a_k \in \mathcal{C} \cup \mathcal{V}$ are the task arguments, $\text{pre}(t_p) \in 2^{\mathcal{P}}$ the preconditions and $\text{eff}(t_p) \in 2^{\mathcal{P}}$ the effects of the primitive task t_p . We further define the positive and negative preconditions of t_p as $\text{pre}^+(t_p) := \{p \in \text{pre}(t_p) : p \text{ is positive}\}$ and $\text{pre}^-(t_p) := \{p \in \text{pre}(t_p) : p \text{ is negative}\}$. We define $\text{eff}^+(t_p)$ and $\text{eff}^-(t_p)$ analogously.

We call a fully ground primitive task an **action**.

As preconditions and effects may not be concerned with the whole world state the closed-world assumption does not apply to them. To any HTN instance we could create an equivalent one where each precondition and effect cares about the whole world state. This would be achieved by instantiating all the "don't care" terms in preconditions and effects with all possible combinations of predicates. Doing this would, however, come at the price of a huge blowup of our planning problem.

Definition 5. An action t_p is **applicable** - pre^+ is fully contained in the state s - none of pre^- are contained in s - the **application** of a task gives us a new state $s' = (s \setminus \text{eff}^-(t_p)) \cup \text{eff}^+(t_p)$

Definition 6. We define a **compound task** as $t_c = \tilde{t}_c(a_1, \dots, a_k)$, where $\tilde{t}_c \in T_c$ is the task symbol from the finite set of compound task symbols T_c and a_1, \dots, a_k are the task arguments.

Primitive and compound tasks together form task networks. In places where both can be used, we will refer to them simply as tasks $t \in T$.

Definition 7. Let $T = T_p \cup T_c$ be a set of primitive and compound tasks. A task network is a tuple $\tau = (T, \psi)$ consisting of tasks T and constraints ψ between those tasks.

Definition 8. Let M be a finite set of method symbols and $T = T_p \cup T_c$ a set of primitive and compound tasks. A **method** $m = (\tilde{m}(a_1, \dots, a_k), t_c, \text{pre}(m), \text{subtasks}(m), \text{constraints}(m))$ is a tuple consisting of the method symbol \tilde{m} , the method arguments a_1, \dots, a_k , the associated compound task $t_c \in T_c$ the method refers to, a set of preconditions $\text{pre}(m) \in 2^P$, a set of tasks $\text{subtasks}(m) = \{t_1, \dots, t_l\}, t_i \in T$ and a set of constraints c_1, \dots, c_m defining relationships between the tasks in $\text{subtasks}(m)$. Any arguments appearing in $t_c, \text{pre}(m), \text{subtasks}(m)$ must also appear in a_1, \dots, a_k .

In TOHTN planning, $\text{constraints}(m)$ is implicitly set s.t. the subtasks t_1, \dots, t_l are totally ordered.

We call a fully ground method a **reduction**.

Each method m has exactly one associated compound task t_c . However, multiple methods m_1, \dots, m_k may be associated with a single compound task t_c . Additionally, while any arguments of t_c must be present in m , the contrary is not true and m may have arguments not present in t_c , i.e., m is not fully determined by t_c . As a result methods present choice points both in the choice of method itself as well as through the argument instantiation.

Definition 9. *Test*

link dire
to resol
on of a
network
resolvin
defined
fore met

2.2 Malleable Algorithms

3 TOHTN Metadata

- researchers in TOHTN planning have collected a number of test instances for TOHTN planning
- provide some analysis of those instances in the context of modelling TOHTN planning as graph search
- provide a foundation to discuss the effects of changes and improvements in the crowd planning framework

4 Implementation

5 Improvements to CrowdHTN

5.1 Reducing Memory Consumption

5.1.1 Efficiently Storing the Preceding Plan

- Use the `ImmutableStack<T>` structure
-

5.1.2 Reducing Copies of the World State

5.1.3 Only Saving 'Potentially Interesting' Nodes

5.2 Making Crowd More CPU Efficient

5.3 Efficiently Hashing Nodes of the Search Graph

Overview:

- the hash of a node consists of two parts
- 1: the hash of the open tasks
- 2: the hash of the world state
- We do not care about the hash of the preceding plan. Nodes with open tasks and world state equivalent have equivalent plans leading to a goal (somewhat similar to the Nerode relation) and we do not care about optimality. How he reach this point with equivalent remaining options is thus not of interest to us
- The hash of the world state can be large, but the number of elements is ultimately bound for any particular instance
- The length of the preceding plan is unbounded

Open tasks:

- Care needs to be taken to use an order independent hash function for the world state, as the underlying representation as a set does not guarantee us any particular iteration order, especially between nodes (might depend on the order of things inserted into the set!)
- Alternatively we could have chosen some other ordered representation, e.g. maintain the world state as a sorted list of predicates with a defined order. This would imply extra work we are unwilling to do.
- For open tasks, we save not only the task itself but additionally the hash of all open tasks from first to current one
- The order in which we hash open tasks is from oldest to newest one
- On applying a Reduction, we push the new open tasks onto the open task stack and compute each of their hashes combined with their predecessors. Each of these hashing operations is completed in $O(1)$
- Effectively, this means that each task is hashed exactly once
- As we already have to push each task onto the open tasks, inserting an additional $O(1)$ operation on pushing does not change the asymptotic runtime

World state:

Time is actually constant world state -> what the expected time?

this is not implemented yet, take care to actually

- In section XXX we discussed how each copy of the world state can be shared by many search nodes to reduce the memory footprint
- Instead of hashing the world state each time on demand, we can store a shared tuple of world state and its hash
- This way, we only hash the world state once, reusing the computation
- Is this actually a time saving?
- Future work: the hash of the world state is order independent (sum of squares of individual hashes), to not worry about forcing any fixed iteration order. Utilize this and wrapping maths to hash only the differential of

5.4 Preceding Plan

In the initial implementation each node stored the full preceding plan as a sequence of all reductions that were applied so far. This leads to roughly quadratic overhead (sum over $1..n$, only roughly as not each step increases the length. Wait, is it roughly, then? Probably, as the fraction of steps that increment the preceding plan should be kinda constant) This duplication was not needed. The newer implementation instead stores an optional<reduction> in each node. I.e., the preceding reduction is stored if one exists, nothing if there isn't one. When the preceding plan is needed, either for communication or to extract a plan, the current search path is iterated and all reductions are accumulated.

5.5 Lazy Instantiation of Child Nodes

lazy instantiation works on the basis of finding all free variables of a method and creating Reductions based on all possible combinations

Initial implementation: instantiate all possible reductions, filter out any with not fulfilled preconditions, then shuffle them

Problems: this spends both time and memory instantiating reductions that might never be needed for the rest of the search We effectively save not only the current path, but also follow all possible branches to a distance of 1

Solution: lazily create reductions as needed How to do this (first way): adapt the argiterator from Lilotane into the CrowdHTN code base. Adapt it to only substitute the arguments that are not already determined by the corresponding task (arguments)

To achieve randomization: each domain is iterated to create the substitutions Each time we build such an iterator, we randomize the order within the domains for this specific iterator This will lead to different orders

Further ideas: each time domains $1..k$ have been fully iterated, increment $k+1$ by 1, then shuffle the order of domains $1..k$

For n total domains, each time domains $1..n-1$ have been iterated, remove the current value from domain n , then shuffle all domain orders

Compare the runtimes of eager and lazy instantiation, check at which point it is worth it to incur the (potential) additional overhead of lazy instantiation Compare on multiple domains? Check different metrics for comparison (size of domains, number of parameters, number of potential children (product of sizes))

Potential problems: We need quite a bit of state (domains, current index into each domain) to perform lazy instantiation The order is not truly random. We iterate some domains faster/

more often than others. What if the important change is in a domain which is iterated slowly? More random order makes this easier

A potential solution: space filling curves Advantages: little state (can just be incremented), iterates all dimensions equally Disadvantage: fixed order. With shuffling within each domain might be random again

Space filling curves come with the restriction of being strictly continuous We do not need this property. All we are interested in is an easy to compute fixed order in which the whole space is iterated where each permutation is hit exactly once We want only self-avoiding curves, to not hit any instantiation twice (could loop detection just fix that? But it would be a bad fix)

5.6 Using Domain Meta-Information

inspiration taken from HyperTension

taking preconditions of tasks and lifting them up into methods Only do this if the precondition is guaranteed to also apply to the method (cannot be achieved before the respective task?) This allows us to stop exploring paths earlier

5.7 Global and Memory Efficient Loop Detection

So far: each worker has a hash set of each node ever encountered (note: we define node equality through the sequence of open tasks and the set of predicates that is the world state. Depth in the graph and preceding plan are ignored) Advantages: allows for perfect local loop detection, as we can always fall back to equality checks in case of hash collisions. Problems: This approach leads to a massive memory overhead, as the world state is duplicated countless times. This also leads to increased run times just to manage the growing hash set. This is also not suited for global loop detection. Global loop detection would need some mechanism to communicate seen nodes. Nodes are too large to communicate all of them, though. If we only communicate some nodes instead of all we run the risk of a node "going past that barrier of known nodes and still exploring the swathes of known nodes beyond the "barrier" TODO: what about open tasks (memory consumption)? Or was that an ImmutableStack with little overhead?

Idea: drop the precondition of perfect loop detection with no false positives. For loop detection use a bloom filter with a set of hash functions for search nodes. This comes with a fixed, configurable memory overhead per worker. As a result, overall memory consumption should - drop - be more predictable To communicate nodes we can just communicate the bit array that makes up our filter and combine them via bitwise OR. This ensures that communication volume is also fixed. It is independent of both size of nodes (i.e. length of open tasks sequence) and number of explored nodes (since last round of communication). In addition to communicating the bloom filter we communicate the number of newly added nodes. Then each worker knows an upper bound for the total number of nodes that are part of it's bloom filter (might be an overestimation if two workers insert identical nodes as the node would be counted twice). As bloom filter performance/ false positives degrade with the number of elements inside it, we can restart our work depending on the number of elements in the bloom filter.

- loop detection information is shared at most every second (or however long it takes for the communication to complete) - as a result, when the root node increases it's version, a broadcast is started that communicates the new version to everyone else - the version in normal messages is still kept. In case a message arrives before the version broadcast, the version is increased even earlier. New nodes likely get their version with the first work package.

5.7.1 Restarts Under Loop Detection

- bloom filters are of fixed size. If we propose a maximum false-positive rate, they will fill up at some point - this necessitates the need for restarts - bloom filter architecture: for local loop detection, use an expanding bloom filter (citation!) - this allows us to not incur any restarts due to local state - for global loop detection, we use a fixed-size bloom filter - if a node is encountered twice for local loop detection, add it to the global loop detector, to communicate - we know that in mallob, the tree of workers is always a 'full' tree, except for the lowest level - i.e. we may be missing some nodes to the right end of the lowest level - especially, the root node always survives - as a result, the root node always receives all the messages about shared loop detection data - i.e., if the global loop detection bloom filter on any node is full, it follows that the global loop detection bloom filter on the root node is also full - the opposite does not hold (i.e. a new node enters just as the root node bloom filter fills up) - as a result, we can simply delegate the question of whether to restart to the root node and have it done in a centralized way, simplifying our communication patterns

5.7.2 Reaching Consensus on Search Version

- each search version corresponds to a restart with subsequent doubling in size of the global loop detector - keeping versions in sync is of importance - we do not want to loose any work - we do not want to insert search nodes from a wrong version into our bloom filter (especially old node into new filter, other way around is discarded either way) - each message between workers is extended by their internal version - if the version of a received message is higher than the internal version, the internal version is increased and then acted accordingly - new workers get updated to the current version the first time they ask for a work package - work requests and responses do not suffice as a version updating mechanism, as work packages may be arbitrarily large and messages arbitrarily rare -

5.7.3 Completeness Under Loop Detection

- any single iteration may be incomplete, e.g. if initial node and goal node hash to the same values - with restarts, we have uniform hashes, compute the chance of colliding each time - this necessitates changing seeds with each restart - soon, we will not collide - yay, plan can be found!

- we do loose termination - so far, termination could be known if the search tree gets too deep (cite limits known from other TOHTN paper) - however, this is infeasible, as RAM is actually finite (refer to known branching factors!) - i.e., we do not know the hash of a goal node, as a result we cannot exclude the possibility of it simply being cut off each time so far - we also do not know the hash of all nodes just before a goal node, etc (the path could always only be a single node wide) - as a result, we cannot terminate at any point, as we cannot exclude the plan hiding in some always-cut-off part of the search space - losing termination is not a big deal anyways, I guess?

6 Theoretical Improvements to the Crowd Planner

6.1 Search Algorithms in Crowd

6.2 Loop Detection

6.2.1 Perfect Loop Detection

6.2.2 Bloom Filters

6.2.3 Distributed Loop Detection

- ensure that

7 Malleable TOHTN Planning

- ensure completeness - ensure performance - work stealing is kinda nice here

8 Experimental Evaluation

9 Future Work

Literatur

- [1] GEORGIEVSKI, ILCHE und MARCO AIELLO: *HTN planning: Overview, comparison, and beyond*. Artificial Intelligence, 222:124–156, 2015.